


RAEv2 Context MLP: recorded training loss

3,000 steps | batch 8 | fixed learning rate 0.0003



Training points are single batches, not 100-step averages; block means use only 5 logged batches.

Only one validation evaluation: final EMA MSE = 0.3052 (32 batches x 8 draws).

No historical validation curve or uncertainty interval is available. Source: retained input_local_training logs.